

What 147.14 tokens/s means in local coding

Qwen3-30B-A3B on this RTX 5070 Ti Laptop GPU (12 GB), kept resident

MEASURED ON THIS LAPTOP - NOT A GUARANTEE FOR EVERY PROMPT

RESIDENT CODING MEASUREMENT

147.14 tokens/s median

90.83 ms median warm cached time to first text

Controlled 256-token average: 149.44 tok/s

IDEALIZED WARM OUTPUT TIME

64 tokens	0.52 s
128 tokens	0.95 s
256 tokens	1.82 s
512 tokens	3.56 s

WHAT THIS FEELS LIKE

**Short coding answers arrive in roughly half a second;
long snippets take a few seconds.**

Across the measured coding requests, 37-225 output tokens took 0.289-1.650 s, with a 0.677 s median. At this speed, reading, testing, and correcting code usually take longer than generation itself.

SPEED WITH MEASURED CODING CORRECTNESS

5/6 task types passed both first runs

6/6 passed after one machine-feedback repair

The repeated failure mutated nested input lists. One repair corrected both runs in 1.18 s median. This is tool-assisted quality, not perfect first-shot accuracy.

LITERAL ACCESS BENEFIT

~58.2 GiB BF16 weights

would need about 4.76x this GPU's total VRAM

- Complete 30.532B-parameter MoE model; 49/49 layers on GPU.
- 8K gate peak: 10,610 MiB; 1,617 MiB headroom.
- External GGUF/IQ2_M carries the low-bit weights; NHDF seals and operates the deployment.

PRACTICAL CONCLUSION

**Generation is fast enough to feel immediate;
correctness checks are now the main delay.**

Keep the server resident to avoid the 12.64 s cold startup. One configured slot means concurrent requests queue.

BOUNDARIES OF THIS FACT

The warm timing used prefix caching. Code has no reliable words-per-token conversion. The accuracy check is six small deterministic Python tasks, not repository-scale coding. Long contexts, thermals, and power state can change throughput; Windows was on Balanced.